

Correcting Step-Cached Diffusion Transformers with a Trajectory Objective

First Author

High School Student, *School Name*
first.author@example.com

Second Author

Affiliation
second.author@example.com

Abstract—Step caching accelerates diffusion transformers by reusing a block’s residual across neighbouring solver steps, which is fast but degrades the image. A learned corrector can predict the reuse error, and the obvious way to train one is to regress the residual it should have produced. We show that this objective is the wrong one, and that replacing it costs twenty minutes of training. The error a cache introduces is *integrated along the sampling trajectory* — a refresh restores the cache, not the latent — so a per-step residual target is defined on a trajectory the corrector never visits. Training the same corrector against the deviation of the latent trajectory instead, by backpropagating through the cached rollout, improves image quality at every injection granularity we tested ($K = 1$ to 28, all significant, $t = 2.7$ to 7.0, $n = 72$). At a fixed step count on PixArt- Σ -XL-2-512, the resulting corrector reaches SSIM 0.6820 against the exact output at $3.41\times$ speedup, beating a reproduction of the TaylorSeer extrapolation mechanism by +0.0946 SSIM ($t = 14.5$, improving on all 72 of 72 paired cases) for 8.4% less speedup, and a reproduction of the Block Caching mechanism by +0.0490 ($t = 8.4$) while being 12% faster than it. We also report the measurement that motivated the change: driving residual accuracy 25–43% better by matched-data retraining buys only +0.005 to +0.013 SSIM, so residual error is a poor guide to what a corrector is worth.

Index Terms—diffusion models, inference acceleration, feature caching, diffusion transformers

I. INTRODUCTION

Sampling from a diffusion transformer costs one full forward pass per solver step. Step caching exploits the fact that consecutive steps produce similar internal features: at an *anchor* step the residual $R = \text{block}(h) - h$ of each transformer block is computed and stored, and for the next $i - 1$ *reuse* steps the block is skipped and its stored residual added instead [1], [2], [3]. This is cheap and training-free, and it visibly degrades the image.

Two families of fixes exist. Training-free extrapolation predicts the residual’s drift from its own history [4], [5], and learned correction trains a small network to predict the reuse error [6], [7]

$$\Delta R = R_{\text{true}} - R_{\text{anchor}}, \quad (1)$$

which is exactly the error verbatim reuse commits. Learned correction is the setting of this paper. Its natural training objective is supervised regression on (1), and every design decision then follows from residual accuracy: architecture, checkpoint selection, and how much correction to inject.

Our contribution is that this objective is misaligned, and that the misalignment is both diagnosable and cheap to fix.

- 1) We measure the deviation of the cached trajectory from the exact one at every solver step, and find that it does *not* reset at cache refreshes: it climbs almost monotonically (0.0089 at step 1 to 0.6092 at step 19). An anchor refreshes the cache, not the latent, so deviation already injected into the state is permanent. Residual error is pointwise; trajectory error integrates.
- 2) We show that pushing residual accuracy hard does not pay. Matching training data volume across granularities and training every rung to convergence improves residual error by 25–43% and image SSIM by only +0.005 to +0.013, against +0.033 for the decision to add correction at all.
- 3) We train the same corrector against the trajectory deviation instead, by backpropagating through the cached rollout. This is inexpensive for a reason specific to caching: a cached block never executes, so the graph over one reuse step is K corrector forwards rather than K transformer blocks. It improves image quality at every K , and by the largest margin where per-site correction degrades worst.
- 4) At a fixed step count it beats reproductions of two published mechanisms on three fidelity metrics, at a speedup cost of 8.4% against one of them and a speedup *advantage* against the other.

II. RELATED WORK

Fewer steps versus cheaper steps: Two orthogonal routes reduce sampling cost. Distillation and consistency training cut the number of steps [8], [9], and high-order solvers do the same without retraining [10]. Caching instead makes each step cheaper, so the two compose. We hold the step count and the solver fixed throughout, which isolates the caching question but also means our comparisons say nothing about which route is preferable — see Limitations.

What to cache: Feature caching for diffusion U-Nets [1] carries over to transformers, where the natural unit is a block’s residual [2], [3], optionally chosen adaptively per step [4] or per token [11]. All of these decide *when* to skip; we take the schedule as given and ask what to add back.

Correcting a cached step: Block Caching [6] applies a timestep-conditioned per-channel scale-shift to the reused feature — trained, but a function of the timestep alone. TaylorSeer [5] forecasts the residual from its own history by

finite differences, requiring no training. Learning to cache [7] learns the skip pattern rather than a correction. Our corrector differs from all three in reading the *current state*, which is what makes error propagation across injection sites multiplicative (Sec. VI-B) and is also why the choice of training objective matters at all: a state-independent correction cannot be improved by observing the trajectory it produces.

Objectives defined on a policy’s own trajectory: Regressing a target collected off-policy and deploying it in closed loop is the setting DAGger addresses in imitation learning [12], and scheduled sampling addresses in sequence models [13]. Both fix the data distribution. Our finding is that in this setting the data distribution is not the binding problem — the objective is — and that because the cached rollout is differentiable, the trajectory can be optimised directly rather than approached by resampling.

III. SETUP

Backbone and schedule: PixArt- Σ -XL-2-512 [14] (0.6B parameters, 28 blocks) with its shipped DPM-Solver++ scheduler [10] at 20 steps and 512² resolution. Anchors are uniform at interval $i = 5$, so 4 anchor steps and 16 reuse steps. All numbers are against the same-seed exact 20-step output.

Corrector: A conditioned MLP (11.9M parameters) reads the state drift $h_t - h_{\text{anchor}}$, the cached residual, and the anchor hidden state, conditioned on the diffusion timestep, the reuse horizon τ , and the site index. At deployment

$$h_{\text{out}} = h_t + R_{\text{anchor}} + \sigma \widehat{\Delta R}, \quad (2)$$

where σ is an injection strength selected on validation and never on test. The state-drift input is what makes the correction state-dependent, which matters in Sec. VI-B. The 28 blocks are grouped into K contiguous cached segments, so K is the number of injection sites per reuse step; $K = 1$ corrects the whole stack once and $K = 28$ corrects every block.

Protocol: Every image-level number is a paired comparison over 24 held-out prompts $\times$ 3 seeds ($n = 72$), with σ re-selected on validation for each corrector over a common grid. We report t statistics and the number of cases that improve. Prompts are templated (10 styles $\times$ 24 subjects $\times$ 10 conditions); the held-out split shares no subject with training.

IV. RESIDUAL ACCURACY IS A POOR GUIDE

Trajectory error integrates: Measuring $\|x_{\text{cached}} - x_{\text{exact}}\|/\|x_{\text{exact}}\|$ at each of the 20 steps shows deviation rising from 0.0089 to 0.6092 with only small dips at the four anchors. This is the basic reason a per-step residual target is misaligned: the corrector’s action at step t changes the state at $t + 1$, while (1) is defined on the unperturbed trajectory.

Better residual accuracy barely helps: To rule out the possibility that the residual objective simply had not been optimised hard enough, we rebuilt training from scratch: data volume matched across granularities at exactly 50,176 slots per rung (a seed multiplier of $28/K$), and the step budget selected per rung on validation by doubling until the gain fell below 0.5%. Residual error improved by 25–43% depending

on K . Image SSIM improved by +0.0131 at $K = 1$ falling to +0.0055 at $K = 7$ — all significant, all small. Two trends run opposite along the granularity axis: the achievable residual improvement rises with K (−27.7% at $K = 1$ to −42.9% at $K = 28$) while the image quality it buys falls.

A ceiling check: Replacing the corrector by an oracle that recovers a fraction β of the true residual gives residual error exactly $(1 - \beta)^2$, so each trained corrector has an equivalent β . Final-step trajectory error falls from 0.6102 at $\beta = 0$ to 0.4384 at $\beta = 0.8$, and to exactly 0 at $\beta = 1$. The curve is convex: the ceiling is perfect and the middle is flat, which rules out the correction form as the limit and locates the problem in the objective. Notably the trained correctors beat an oracle at their own equivalent β by about 1.5 $\times$, because a scalar oracle shrinks error isotropically while a trained corrector learns which directions matter to the trajectory. Residual error therefore distorts in both directions: as an objective it overstates the gain, as a capability measure it understates the corrector.

V. TRAINING AGAINST THE TRAJECTORY

We keep the architecture, the schedule and the deployment rule (2), and change only the loss. Let $x_1, \dots, x_T$ be the cached latents and $x_1^*, \dots, x_T^*$ the exact ones under the same seed. The objective is

$$\mathcal{L} = \sum_{t \in \text{reuse}} \frac{\|x_t - x_t^*\|^2}{\|x_t^*\|^2}, \quad (3)$$

minimised by backpropagating through the cached rollout. Supervision is free: the exact trajectory comes from the same forward pass that produces the corrector’s input, so no annotation is involved.

Two implementation facts make this practical, and both are specific to caching. First, a cached block never runs its transformer block — the runtime returns $h_t + R_{\text{anchor}} + \sigma \widehat{\Delta R}$ — so the graph over one reuse step consists of K corrector forwards, not K blocks of a 0.6B model. Peak memory is 7.5 GB at $K = 1$; checkpointing each corrector call brings $K = 14$ from over 44 GB down to 4.2 GB. Second, gradients are cut at anchor steps, where the cache is refreshed from exact values, so a corrector call can only influence later intervals through the latent. Extending the graph across intervals improves (3) but does not improve image metrics, so we keep the one-interval truncation. Training is 400 steps, about 20 minutes on one A40.

VI. RESULTS

A. Fixed step count against published mechanisms

Table I compares at 20 steps and $i = 5$, so every arm skips the same blocks and speedup differs only by each method’s correction overhead. We reproduce the TaylorSeer mechanism [5] (Newton divided-difference forecast of the residual from its history, training-free) and the Block Caching mechanism [6] (a timestep-conditioned per-channel scale-shift, trained but state-independent). Both are reimplemented in our runtime rather than run from the authors’ code, so they are reproductions

TABLE I
FIXED 20 STEPS, $i = 5$, $n = 72$. “Sp.” IS MEASURED WALL-CLOCK SPEEDUP AGAINST THE EXACT MODEL; Δ SSIM IS PAIRED AGAINST VERBATIM CACHING. THE TWO REPRODUCED MECHANISMS ARE DESCRIBED IN THE TEXT.

Method	Sp.	SSIM $\uparrow$	PSNR $\uparrow$	LPIPS $\downarrow$	Δ SSIM
Exact, 20 steps	1.72	0.7951	21.78	0.1786	—
Verbatim cache	3.79	0.6269	15.37	0.4669	—
TaylorSeer, 1st	3.72	0.5874	14.31	0.5080	−0.0395
TaylorSeer, 2nd	3.73	0.4782	13.57	0.6386	−0.1487
Block Caching	3.04	0.6330	15.66	0.4594	+0.0060
Residual obj.	3.41	0.6676	16.48	0.3795	+0.0407
Trajectory obj.	3.41	0.6820	17.20	0.3883	+0.0551




Fig. 1. Left: gain over verbatim caching at each granularity under the two objectives, each at its own validation-selected σ . Right: sensitivity to σ for the high- K correctors; the optimum moves left and the peak narrows as K grows, and $K = 14$ and $K = 28$ turn negative by $\sigma = 0.25$.

of the mechanisms under our schedule and backbone, not replications of the published results.

Paired against the trajectory-objective corrector, the first-order TaylorSeer mechanism loses 0.0946 SSIM ($t = 14.5$) and 0.1197 LPIPS ($t = 14.2$), with our corrector better on all 72 of 72 cases; the second-order version loses 0.2038 ($t = 39.8$, 72/72); Block Caching loses 0.0490 ($t = 8.4$, 64/72). The exchange rate is therefore favourable: 8.4% less speedup than TaylorSeer for +0.0946 SSIM, and against Block Caching we are both faster and better.

Two qualifications belong with the table. TaylorSeer at $i = 5$ is *worse than no correction at all* (−0.0395, only 3 of 72 cases improve); in our reproduction it helps at $i = 2$ –3 and is past its useful extrapolation range by $i = 5$, so Table I is not evidence against it at its own operating point. And TaylorSeer is training-free, while our corrector costs training time, 11.9M parameters and 0.041 TFLOP per correction — which is exactly where the 8.4% goes.

B. The objective change helps at every granularity

Table II and Fig. 1 show the objective change helping at all six granularities, every difference significant. The margin is largest where per-site correction degrades worst: +0.0370 at $K = 7$ against +0.0147 at $K = 1$. This is consistent with the mechanism behind the degradation. A state-dependent correction turns error propagation across sites multiplicative rather than additive — $\delta_{b+1} \approx (I + J_b)\delta_b + \epsilon$ with $J_b = \partial\Delta R/\partial h$ — and the trajectory objective is the only one that observes the consequence of that compounding rather than

TABLE II
PER-GRANULARITY COMPARISON OF THE TWO OBJECTIVES, $n = 72$, EACH AT ITS OWN VALIDATION-SELECTED σ .

K	Residual obj.	Trajectory obj.	Difference
1	+0.0406	+0.0553 ($\sigma=0.60$)	+0.0147 ($t = 2.7$)
2	+0.0149	+0.0385 ($\sigma=0.30$)	+0.0236 ($t = 3.8$)
4	+0.0189	+0.0450 ($\sigma=0.22$)	+0.0261 ($t = 5.0$)
7	+0.0145	+0.0515 ($\sigma=0.18$)	+0.0370 ($t = 7.0$)
14	+0.0096	+0.0414 ($\sigma=0.18$)	+0.0318 ($t = 4.7$)
28	+0.0148	+0.0380 ($\sigma=0.18$)	+0.0232 ($t = 4.8$)




Fig. 2. Native-resolution 250² crops, same seed and prompt across a row, SSIM against the exact output below each panel. Crops rather than downsampled images because thumbnails hide the texture differences these methods differ in.

pointwise error. The degradation is weakened, not removed: the residual objective’s gain falls 3–4 $\times$ from $K = 1$ to high K while the trajectory objective’s falls 1.2–1.8 $\times$.

Fig. 1 (right) also shows why σ must be re-selected per corrector. The optimum moves from 0.60 at $K = 1$ to 0.18 at high K , and the peak narrows sharply: $K = 4$ stays above +0.039 across $\sigma \in [0.18, 0.25]$ while $K = 14$ and $K = 28$ are already negative at $\sigma = 0.25$. Both corrector generations trace the same family of curves, so this is a property of granularity rather than of the objective.

Fig. 2 shows the same three prompts under each method, and the failure modes differ in kind rather than in degree. Verbatim caching smears surface detail into streaks and flattens the background gradient. The TaylorSeer mechanism’s dominant artefact is not lost detail but *colour* — extrapolating the residual introduces chromatic fringing along edges, which is consistent with it scoring worst on LPIPS in Table I despite skipping the same blocks. The Block Caching mechanism keeps structure cleaner but renders tone flatter and leaves local bloom artefacts, which its per-channel scale-shift can shift but not localise. Our correction recovers the background gradient most closely and leaves the least fringing, though not none.

VII. LIMITATIONS

One backbone, and it saturates early: Measuring the exact model against step count reveals that ImageReward [15] is saturated by 10 steps on this backbone (102% of a 100-step reference) while pixel fidelity keeps improving. Any method selling saved compute therefore competes against an already-saturated baseline here, and an equal-compute comparison on this backbone favours simply lowering the step count. That is a property of PixArt- Σ with a high-order solver, not a verdict on caching, which is why we compare at fixed step count. Whether our results transfer to a backbone that needs many steps is untested.

Metrics disagree: We report SSIM, PSNR and LPIPS because no single one is decisive. An objective trained on LPIPS wins on LPIPS (+0.1067 over verbatim caching) and loses on SSIM and PSNR; ImageReward cannot separate the objectives at all ($t = 0.2$ to 1.2), though it does separate every corrector from verbatim caching (+0.33 to +0.40, all $t > 4$). Which reward to prefer is not answerable from our evidence.

Scope: Templated prompts (24 training subjects), 512^2 only, 24 held-out prompts $\times$ 3 seeds, and a corrector trained for 400 steps without hyperparameter search from a residual-trained initialisation — so the reported gain is marginal over that initialisation rather than a standalone method. Correctors trained to low residual error can overflow fp16 at deployment; we deploy the high- K ones in fp32.

VIII. CONCLUSION

The natural objective for a learned cache corrector — regress the residual it should have produced — is misaligned with what caching actually damages, because cache error integrates along the sampling trajectory while residual error is pointwise. Optimising the residual objective harder is a poor investment: a 25–43% residual improvement buys +0.005 to +0.013 SSIM. Replacing the objective with the trajectory deviation, which costs 20 minutes of training and no architectural change, improves image quality at every injection granularity we tested, and at a fixed step count beats reproductions of two published mechanisms on three fidelity metrics.

REPRODUCIBILITY

Code, configurations, per-run logs and the raw per-case metrics behind every number are available at <https://github.com/xiaruize0911/Diff-cache-Comp>. Feature collection is seeded and reproducible bit-for-bit; every table cell can be recomputed from the archived `results.json` files.

REFERENCES

- [1] X. Ma, G. Fang, and X. Wang, “Deepcache: Accelerating diffusion models for free,” in *CVPR*, 2024.
- [2] P. Selvaraju, T. Ding, T. Chen, I. Zharkov, and L. Liang, “Fora: Fast-forward caching in diffusion transformer acceleration,” *arXiv preprint arXiv:2407.01425*, 2024.
- [3] P. Chen, M. Shen, P. Ye, J. Cao, C. Tu, C.-S. Bouganis, Y. Zhao, and T. Chen, “ δ -dit: A training-free acceleration method tailored for diffusion transformers,” *arXiv preprint arXiv:2406.01125*, 2024.
- [4] F. Liu, S. Zhang, X. Wang, Y. Wei, H. Qiu, Y. Zhao, Y. Zhang, Q. Ye, and F. Wan, “Timestep embedding tells: It’s time to cache for video diffusion model,” *arXiv preprint arXiv:2411.19108*, 2024.
- [5] Liu *et al.*, “From reusing to forecasting: Accelerating diffusion models with taylorseers,” in *ICCV*, 2025, arXiv:2503.06923.
- [6] F. Wimbauer, B. Wu, E. Schoenfeld, X. Dai, J. Hou, Z. He, A. Sanakoyeu, P. Zhang, S. Tsai, J. Kohler *et al.*, “Cache me if you can: Accelerating diffusion models through block caching,” in *CVPR*, 2024.
- [7] X. Ma, G. Fang, M. B. Mi, and X. Wang, “Learning-to-cache: Accelerating diffusion transformer via layer caching,” in *NeurIPS*, 2024.
- [8] T. Salimans and J. Ho, “Progressive distillation for fast sampling of diffusion models,” *ICLR*, 2022.
- [9] Y. Song, P. Dhariwal, M. Chen, and I. Sutskever, “Consistency models,” in *ICML*, 2023.
- [10] C. Lu, Y. Zhou, F. Bao, J. Chen, C. Li, and J. Zhu, “Dpm-solver: A fast ode solver for diffusion probabilistic model sampling in around 10 steps,” *NeurIPS*, 2022.
- [11] C. Zou, X. Liu, T. Liu, S. Huang, and L. Zhang, “Accelerating diffusion transformers with token-wise feature caching,” *arXiv preprint arXiv:2410.05317*, 2024.
- [12] S. Ross, G. J. Gordon, and D. Bagnell, “A reduction of imitation learning and structured prediction to no-regret online learning,” in *AISTATS*, 2011.
- [13] S. Bengio, O. Vinyals, N. Jaitly, and N. Shazeer, “Scheduled sampling for sequence prediction with recurrent neural networks,” in *NeurIPS*, 2015.
- [14] J. Chen, C. Ge, E. Xie, Y. Wu, L. Yao, X. Ren, Z. Wang, P. Luo, H. Lu, and Z. Li, “Pixart- σ : Weak-to-strong training of diffusion transformer for 4k text-to-image generation,” in *ECCV*, 2024.
- [15] J. Xu, X. Liu, Y. Wu, Y. Tong, Q. Li, M. Ding, J. Tang, and Y. Dong, “ImageReward: Learning and evaluating human preferences for text-to-image generation,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.